






DELIVERED WORK — the same runtime, four domains, judged by tools and people it does not control

AI for INFRASTRUCTURE GPU kernels · serving · local

flash-linear-attention 4 kernels
RWKV6 1.21× · KDA 1.29× · AttnRes 1.10× · SM100 fix
reviewed by outside maintainers; no blocking changes

FlashDA — dLLM masks on FA4 19/19
6 mask families, SM80+SM90; within 5–14% of native FA4
2.5–4.2× over Triton · under ¥80 of model spend

SGLang · MiniMax-H3 +11,263 ln
unified model integrated: 160/160 exact, 5.108× at BS8
62 GiB video model on a 24 GB laptop · 6.159× on A6000

AI for SCIENCE materials · mathematics

MOF chemical control 92.5 / 100 / 74.5%
metal · node nuclearity · ligand family, from 3,300 requests
only 16.7% of parameters trained; 4/8×B200 run by Argus

Structure judgment — AUC 0.594 → 0.833
a simpler sampler beat the published one at equal compute,
admitted only after removing a confound from the comparison

Erdős–Gyárfás campaign 6 updates
six proof-backed frontier updates; one falsified route kept
with its evidence, so later work need not re-enter it

AI for SILICON ACE-2 accelerator

Functional closure 13,914 / 13,914
24-layer Qwen2.5-0.5B W4A8, 18 operators, bit-exact
certified by a Reviewer that did not do the work

Mapped SKY130 synthesis 0.614 mm²
62,283 cells at 100 MHz; WNS and TNS both 0.00 ns
the certificate publishes its exclusions, not just passes

Human involvement 8 days · 6 asks
roughly one interruption every 32 hours; 2 of the 6
were the server dying, not the runtime asking