

Six completed manuscripts - real PDF first pages

Scientific result on each card - Argus production process below

From Leaderboard Ranks to Audited Decisions: A Controlled-Panel Study of Evaluation-Format Uncertainty

Anonymous ACL submission

[illegible]

1 Introduction

When Structure Does Not Help: A Diagnostic Study of Training-Free Gates for Compositional Image-Text Matching

Anonymous ACL submission

Abstract

Computational models are becoming much more sophisticated, and the number of models that can be used to simulate a particular system is increasing. This paper discusses the challenges that arise from the fact that many simulations are right now being done on a single machine, and the challenges that arise from the fact that many simulations are now being done on a distributed system. The paper discusses the challenges that arise from the fact that many simulations are now being done on a distributed system, and the challenges that arise from the fact that many simulations are now being done on a distributed system. The paper discusses the challenges that arise from the fact that many simulations are now being done on a distributed system, and the challenges that arise from the fact that many simulations are now being done on a distributed system.

1 Introduction

Computational models are becoming a de facto standard in global system analysis. The paper discusses the challenges that arise from the fact that many simulations are now being done on a distributed system, and the challenges that arise from the fact that many simulations are now being done on a distributed system. The paper discusses the challenges that arise from the fact that many simulations are now being done on a distributed system, and the challenges that arise from the fact that many simulations are now being done on a distributed system.

When a Frozen Gate Refuses to Adapt: A Diagnostic Audit of Test-Time Adaptation in Vision-Language Models

Anonymous submission

[illegible]

Introduction

Language models such as CLIP provide strong zero-shot generalization (Radford et al., 2021; Bhoja et al., 2021). Deployment-ready neural networks can be used to generate embeddings for images in various, open, commonly used scenarios (Radford et al., 2021; Bhoja et al., 2021). Existing models can be used to generate embeddings for images in various, open, commonly used scenarios (Radford et al., 2021; Bhoja et al., 2021). Existing models can be used to generate embeddings for images in various, open, commonly used scenarios (Radford et al., 2021; Bhoja et al., 2021).

Surface-Prior Selection for Training-Free GUI Reliability at Matched Coverage

Anonymous ACL submission

[illegible]

Introduction

GUIs increasingly require attention on automation, mobile computing, and page-to-action data interfaces. This shift is visible in recently developed prompting models such as ScreenCraft and GPT4o (OpenAI, 2024), which are designed to revolutionize professional GUI benchmarks such as ScreenPilot-Pro (Li et al., 2025), coordinate-free GUIs for automating forms (e.g., GUC4, Action and OmniParser (Wu et al., 2025; Lu et al., 2024)), and broader GUI understanding models (e.g., ScreenAI and Forms 4U (Bachtler et al., 2024; Bachtler et al., 2025)).

This paper adds a new perspective: under ideal conditions, can a training-free wrapper around a frozen GUI evaluator predict real high-level control wiring (e.g., the sequence of GUI elements to click) for a GUI? The evaluated system uses control (Qwen2.5-VL (Qwen Team, 2025)), it trains no model, uses no data, and is trained on a single dataset (Qwen2.5-VL (Qwen Team, 2025)). Matched coverage is the central methodological device. Within each benchmark, we compare the coverage of the wrapper against the number of elements, the number of actions it executes the same number of rows, so abstraction is self-contained to the hidden information. The method uses action-coat labels as the main ordering signal

When Frozen Multimodal Hallucination Wrappers Fail

Anonymous submission

[illegible]

This paper asks a deliberately naive question: under a retained public-landmark status, when do calibration and/or validation tests need to be repeated? The answer to this question is not obvious, and the authors present a series of simulations to explore this question. The authors conclude that the need to repeat calibration and/or validation tests depends on the number of landmarks used, the number of landmarks used during time, attention, and displacement, and the number of landmarks used in the validation tests. The authors conclude that the need to repeat calibration and/or validation tests depends on the number of landmarks used, the number of landmarks used during time, attention, and displacement, and the number of landmarks used in the validation tests. The authors conclude that the need to repeat calibration and/or validation tests depends on the number of landmarks used, the number of landmarks used during time, attention, and displacement, and the number of landmarks used in the validation tests.

When Vocabulary-Row Proxies Fail: A Negative Result for Calibration-Free Vocabulary Quantization

Anonymous ACL submission

[illegible]

Evaluation reliability

Audited leaderboard decisions

15 stable - 300 abstained

Vision-language matching

Compositional matching gate

0.953 fallback > 0.942
gate

Test-time adaptation

Frozen adaptation gate

2 / 1,880 DTD updates

Multimodal agents

GUI reliability at matched coverage

0.0125 -> 0.0115 risk area

Multimodal hallucination

Frozen hallucination wrappers

15 cells - 4,500 rows

Model quantization

Vocabulary-row quantization proxies

0 significant wins vs random

ARGUS

production process

6/6
papers completed

640
campaign-hours

576

Engineer rounds

286
*Reviewer revisions*89
session rolls

16 Stage rollbacks